



# OpenAI Agents SDK Documentation Overview

This report consolidates information from the **OpenAI Agents SDK for TypeScript** documentation into a single reference. It covers all guides and extensions from the site (<https://openai.github.io/openai-agents-js/>), summarizing concepts, code snippets and features. Citations link back to the original pages.

## Overview

The OpenAI Agents SDK helps developers build **agentic AI applications** in TypeScript/JavaScript. It provides primitives such as **Agents**, **Tools** and **Handoffs**, enabling complex workflows with built-in tracing, guardrails, sessions and support for both text and real-time voice interactions <sup>1</sup>. Key features include:

- **Agent loop** – the SDK implements an iterative loop where the LLM is called, the response is inspected for final output, tool calls or handoffs, and the loop repeats until a stopping condition (final answer, error or max turns) <sup>2</sup>.
- **TypeScript-first** – strong typings via Zod/JSON schemas for tool inputs and agent outputs <sup>3</sup>.
- **Guardrails** – input/output/tool validations to enforce safety rules <sup>4</sup>.
- **Sessions** – persistent memory layer to preserve conversation history across runs <sup>5</sup>.
- **Tracing** – automatic collection of spans for debugging and performance analysis <sup>6</sup>.
- **Real-time voice** – integration with OpenAI's speech-to-speech models enabling streaming audio chat <sup>7</sup>.

Installation is via `npm install @openai/agents` (with `zod` dependency). A simple “Hello world” agent can be created by defining instructions and calling `run(agent, input)` <sup>1</sup>.

---

## Quickstart

The quickstart guide walks through setting up a project:

1. **Initialize project** – use `npm create vite@latest` to scaffold a TypeScript project and install `@openai/agents` & `zod` `[2443362791021±L872-L1152]`.
2. **Set API key** – set `OPENAI_API_KEY` in your environment or call `setDefaultOpenAIKey()` <sup>8</sup>.
3. **Create an agent** – instantiate `Agent` with a `name` and `instructions`. Example:

```
const agent = new Agent({
  name: 'Writer',
  instructions: 'Write a haiku about cats.'
```

```
});  
const answer = await run(agent, 'Compose a haiku');
```

1. **Add tools** – define function tools using `tool()` with Zod schemas for parameters; attach them to the agent. For example, a weather tool returns weather for a city [9](#).
2. **Multi-agent orchestration** – create specialized agents (e.g., billing/refunds) and a triage agent that chooses the appropriate agent using handoffs. Use `Runner.run()` to orchestrate and inspect `result.output` and `result.history` [\[2443362791021±L872-L1152\]](#).
3. **View traces** – the OpenAI dashboard can display traces recorded during the run for debugging and optimization [\[2443362791021±L872-L1152\]](#).

## Agents

An **Agent** represents an LLM configured with instructions, a model, tools and optional handoffs. Key aspects:

- **Configuration** – an agent is created with options: `name`, `instructions`, `prompt / promptId`, `handoffDescription`, `model` (default is `gpt-4.1`), `modelSettings` (temperature, topP, etc.), `tools`, `handoffs` and connectors such as MCP servers [3](#).
- **Context** – two types of context exist: *local context* (dependency injection passed to tools via `RunContext`) and *agent/LLM context* (information visible to the LLM). Local context is not exposed to the LLM and can hold secrets or functions [10](#).
- **Output types** – without a schema, agents return strings; with a Zod schema or JSON schema the output is structured; handoffs allow union output types [11](#).
- **Multi-agent patterns** – agents can hand off sub-tasks to specialized agents (`transfer_to_<agent_name>` tool). Patterns include manager vs. specialist, sequential chains, parallel agents, critique-evaluate loops and code orchestrations [12](#).

## Running Agents

Agents are executed via `run()` or by constructing a `Runner`. Running an agent performs the loop: call the model, inspect response, process tool calls/handoffs, then continue until final output. Options include:

- **Streaming** – set `{ stream: true }` to get an async iterator of events. Use `toTextStream()` to extract only text. Interruptions (e.g., tool approval) can be resolved and the run resumed [13](#).
- **Run arguments** – set `maxTurns`, `context`, `session` (for memory), `stream`, `tracing`, `errorHandlers`, `conversationId` and more [2](#).
- **Run results** – `RunResult` holds `output` (final answer), `history` (inputs/outputs), `lastAgent`, `newItems` and `interruptions` (pending tool approvals) [14](#). `StreamedRunResult` is an async iterator over events.

## Results

The **Results** guide explains how to handle run outputs:

- Final output may be a string or typed object depending on agent configuration and schemas <sup>11</sup>.
  - Access history and new items via `result.history` and `result.newItems` to continue the conversation or resume after interruptions <sup>14</sup>.
  - `interruptions` array contains pending approvals for human-in-the-loop scenarios; use `state.approve()` or `state.reject()` to continue <sup>14</sup>.
- 

## Tools

Tools allow agents to perform actions. Categories:

1. **Hosted OpenAI tools** – web search, file search, code interpreter and image generation (provided via helper functions) <sup>15</sup>.
2. **Local built-in tools** – `computerTool`, `shellTool`, `applyPatchTool` enable interacting with the user's environment or applying patches to files <sup>16</sup>.
3. **Function tools** – wrap arbitrary functions with JSON schema or Zod for structured input; pass options like `needsApproval` or `isEnabled` <sup>9</sup>.
4. **Agents as tools** – expose another agent as a tool using `agent.asTool()`.
5. **MCP servers** – integrate remote tools via Model Context Protocol; connect hosted servers or local servers for retrieval and tool execution <sup>17</sup>.
6. **Codex (experimental)** – call code generation models.

Each tool returns results via messages in the run history. Tools can be disabled or require human approval before execution.

---

## Orchestrating Multiple Agents

The SDK supports two orchestration approaches:

- **Via the LLM** – supply a manager agent that decides which sub-agent to call using handoffs and tools. Provide clear instructions and recommended prompts; monitor responses and introspect reasoning <sup>12</sup>.
  - **Via code** – deterministically chain agents using structured outputs or loops; run agents in sequence or parallel; use a critique-evaluate loop where one agent critiques another's output; or process tasks in parallel to improve performance <sup>12</sup>.
-

## Handoffs

Handoffs allow an agent to delegate part of a conversation to another agent. To define handoffs:

- Specify `handoffs` array in the agent configuration or use `handoff()` helper to register a sub-agent <sup>18</sup>.
  - Each handoff creates a tool named `transfer_to_<agent_name>`; customizing the name and description is possible. Input types can be defined to control what data is passed during handoff <sup>18</sup>.
  - Use callbacks (`onHandoff`) for custom behaviors or to override default prompts and input filters.
- 

## Context Management

Context appears in two forms:

- **Local context** – an object passed to `run()` that holds dependencies or data for tools (e.g., user info, database connections). Tools access it via the `RunContext` argument. This context is *not* visible to the language model <sup>10</sup>.
  - **Agent/LLM context** – content visible to the LLM, provided via instructions, input to `run()`, retrieval tools or webs. Additional context can be added by retrieving data and injecting it into the LLM prompt <sup>10</sup>.
- 

## Sessions

Sessions provide persistent memory across runs. The SDK offers:

- **MemorySession** – in-memory session for simple cases.
- **OpenAIConversationsSession** – connects to the Conversations API; supply options such as `conversationId`, `client`, `apiKey`, `baseURL` and `organization` <sup>19</sup>.

A `Session` implementation stores conversation history; when passed to `Runner.run()` the runner automatically fetches previous items and persists new ones <sup>5</sup>. Sessions can be customized by implementing the `Session` interface.

---

## Models

Agents call language models via the `Model` interface. Important points:

- **Default model** – if not specified, the SDK uses `gpt-4.1`. You can override via agent config, set `OPENAI_DEFAULT_MODEL` environment variable or define default model on `Runner` <sup>20</sup>.
- **Model settings** – `ModelSettings` parameters (temperature, topP, frequency/presence penalties, tool choice, etc.) can be set per agent or run <sup>21</sup>.

- **Prompt reuse** – pass `promptId` to use a server-stored prompt with variables; this allows centralizing prompt management <sup>22</sup>.
  - **Custom model providers** – implement `ModelProvider` and `Model` classes to wrap non-OpenAI LLMs; example `EchoProvider` echoes input and can be plugged into `Runner` [627821135777929±L1218-L1303].
  - **Tracing exporter** – export traces by setting tracing API key using `setTracingExportApiKey()` <sup>23</sup>.
- 

## Guardrails

Guardrails validate inputs, outputs or tool calls:

- **Input guardrails** – run before the model call; receive the user input and return `tripwireTriggered` and optional replacement text. If triggered, execution stops <sup>24</sup>.
- **Output guardrails** – run on final agent responses to check for policy violations and optionally replace or block the output <sup>25</sup>.
- **Tool guardrails** – wrap function tools to validate input/output; may allow, modify, reject or throw errors <sup>26</sup>.

Guardrails can run in parallel or sequentially; sequential mode reduces token cost but increases latency <sup>27</sup>. Guardrail examples include blocking secrets, redacting PII, or verifying that tool output is appropriate [599932972234342±L959-L1258].

---

## Streaming

Enabling streaming returns an async iterator of events for real-time updates:

- **Usage** – call `run(agent, input, { stream: true })` to receive a `StreamedRunResult` with `.toTextStream()` and `.completed` promise <sup>13</sup>.
  - **Event types** – events include `raw_model_stream_event` (model deltas), `run_item_stream_event` (run items like tool calls), and `agent_updated_stream_event` (agent updates) <sup>28</sup>.
  - **Human-in-the-loop** – during streaming, `stream.interruptions` lists pending tool approvals; call `state.approve()` / `state.reject()` and resume by re-running with `{ stream: true }` <sup>29</sup>.
  - **Tips** – always await `stream.completed`, re-specify `stream` when resuming, and use `toTextStream()` to get only text <sup>30</sup>.
-

## Human-in-the-Loop

To require user approval for tool calls:

- Set `needsApproval` to `true` or provide a function on the tool definition to dynamically decide if approval is required <sup>31</sup>.
  - When an agent attempts to call such a tool, execution pauses and `RunResult.interruptions` contains a `ToolApprovalItem`. The developer must approve or reject via `state.approve(item)` or `state.reject(item)` before resuming <sup>32</sup>.
  - In long approval scenarios, serialize `result.state` (e.g., `JSON.stringify(result.state)`) and later reconstruct with `RunState.fromString()` to resume <sup>33</sup>.
- 

## Model Context Protocol (MCP)

The **Model Context Protocol** integrates external tool servers. Three server types are supported <sup>17</sup>:

1. **Hosted MCP servers** – remote servers used by OpenAI's Responses API. Connect using `hostedMcpTool()` and configure approval policies for tool calls. Use connectors to integrate with third-party services like Google Calendar by specifying `connectorId` and `authorization` <sup>34</sup>.
2. **Streamable HTTP servers** – build a server exposing tools and retrieval; connect via `MCP Sever StreamableHttp` with server URL <sup>35</sup>.
3. **Stdio servers** – run tools via a local command (e.g., `npx` to run a server) using `MCP Sever Stdio`. You can define the command, encoding and environment <sup>36</sup>.

Use `connectMcpServers()` to connect multiple servers; it returns lists of active and failed servers <sup>37</sup>.

---

## Tracing

Tracing collects spans for agent runs and is enabled by default on server environments <sup>38</sup>. Important points:

- **Export** – in most environments, traces are exported automatically; in browsers or Cloudflare Workers, call `getGlobalTraceProvider().forceFlush()` to manually flush <sup>39</sup>.
  - **Spans vs traces** – a trace represents one workflow; spans represent operations such as agent runs, model generations, tool calls, guardrails or handoffs <sup>40</sup>.
  - **Default tracing** – the SDK wraps `run()` and other operations into spans. Use `RunConfig.workflowName` or `withTrace()` to override the default name <sup>6</sup>.
  - **Custom traces** – use `withTrace()` or `getGlobalTraceProvider().createTrace()` to create nested traces and spans; add custom processors via `addTraceProcessor()` <sup>41</sup>.
-

## Configuration & Troubleshooting

- **API keys** – set `OPENAI_API_KEY` environment variable, call `setDefaultOpenAIKey()`, or customize the OpenAI client using `setDefaultOpenAIclient()` <sup>42</sup>. Switch between Chat Completions and Responses API with `setOpenAIAPI()` <sup>42</sup>.
  - **Tracing settings** – enable or disable tracing via `setTracingExportApiKey()` and `setTracingDisabled(true)` <sup>43</sup>.
  - **Debug logging** – enable verbose logs by setting `DEBUG=openai-agents*`; additional variables control whether model/tool data is logged <sup>44</sup>.
  - **Supported environments** – Node.js 22+, Deno, Bun are fully supported; Cloudflare Workers and browsers have limitations (manual flush for traces, restrictions on web sockets) <sup>45</sup>. Use debug logs to troubleshoot issues <sup>46</sup>.
  - **Versioning** – the package uses `0.Y.Z` semantics; Y increments indicate breaking changes, Z increments are new features/bug fixes <sup>47</sup>.
- 

## Voice Agents

Voice Agents leverage OpenAI's speech-to-speech models for **real-time voice chat**. They support streaming audio, text and tool calls and are suitable for phone support, mobile apps or voice chat <sup>7</sup>. The Voice Agents SDK provides a TypeScript client for the **OpenAI Realtime API** <sup>7</sup>.

### Key Features <sup>48</sup>

- Connect over WebSocket or WebRTC.
- Works in both browser and backend environments.
- Supports audio and interruption handling.
- Enables multi-agent orchestration through handoffs.
- Defines and calls tools (function tools and hosted MCP tools).
- Supports custom guardrails to monitor model output.
- Provides callbacks for streamed events.
- Reuses components from text agents for voice agents.

### Quickstart <sup>49</sup>

1. **Set up project** – create a Vite TypeScript app and install `@openai/agents` and `zod`. For browser-only usage, install `@openai/agents-realtime`.
2. **Generate client token** – since voice agents run in the browser, obtain an **ephemeral client key** from the Realtime API (starts with `ek_`). Use your OpenAI API key to request a client secret via `curl` <sup>50</sup>.
3. **Create a `RealtimeAgent`** – similar to a regular agent; specify name and instructions <sup>51</sup>.
4. **Create a `RealtimeSession`** – pass the agent and model (e.g., `gpt-realtime`). The session handles audio processing, interruptions and lifecycle <sup>52</sup>.
5. **Connect to the session** – call `session.connect({ apiKey: 'ek...' })` to connect via WebRTC (browser) or WebSocket (server) <sup>53</sup>.

6. **Combine code** – import `RealtimeAgent` and `RealtimeSession`, create agent and session, and call `session.connect()` with the key. Grant microphone access; you can then speak to the agent [54](#).

## Building Voice Agents

The [Building Voice Agents](#) guide dives deeper:

### Session Configuration [55](#)

- Pass additional options when constructing `RealtimeSession` or calling `connect()`. Options include `model`, audio formats (`inputAudioFormat`, `outputAudioFormat`), transcription settings (`inputAudioTranscription.model`), local context, whether to store audio in history, output guardrails, tracing settings, group IDs, custom metadata, auto-trigger options and custom error formatters [55](#).

### Handoffs [56](#)

- Voice agents can delegate tasks to other `RealtimeAgent`s similar to text agents. Handoffs update the ongoing session with the new agent configuration; input filters aren't applied and the model/voice cannot be changed during a handoff [57](#).

### Tools [58](#)

- Voice agents support **function tools** (executed locally) and **hosted MCP tools** (executed remotely). Define tools with `tool()` and attach them to a `RealtimeAgent`. Function tools run in the same environment (browser or server); sensitive actions should call a backend API [58](#). Use `backgroundResult()` to return tool results without immediately triggering a model response [59](#). Tool timeouts behave similarly to text agents [60](#).

### Conversation History [61](#)

- Tools can access a snapshot of conversation history via the `details.context.history` parameter; note the last user phrase may not yet be transcribed [61](#).

### Approval Before Execution [62](#)

- Setting `needsApproval: true` on a tool triggers a `tool_approval_requested` event. The developer should present a UI to approve or reject the call using `session.approve()` or `session.reject()` [62](#).

### Guardrails [63](#)

- **RealtimeOutputGuardrail** functions run asynchronously during voice output to cut off responses that violate rules. They emit `guardrail_trippped` events. Guardrails run every 100 characters by default; adjust with `outputGuardrailSettings` [63](#).

## Turn Detection & Interruptions 64

- Voice activity detection determines when the user starts/stops speaking; configure via `turnDetection` with settings like `type`, `eagerness`, `createResponse` and `interruptResponse` 65.
- Interruptions occur when a user talks over the agent; the session emits `audio_interrupted`. Developers can also manually call `session.interrupt()` to stop audio playback 64.

## Text Input & History Management 66

- Send text via `session.sendMessage(text)` to provide additional context or enable multi-modal input 67.
- `session.history` stores the conversation; listen for `history_updated` events and use `session.updateHistory()` to modify or remove messages 68.

## Limitations & Delegation 69

- You cannot update function tool calls after execution; text output requires transcripts; truncated responses lack transcripts 70.
- Delegation through tools allows combining history and a tool call to offload complex tasks to another backend agent and return results to the user 71.

## Transport Mechanisms

The **Realtime Transport Layer** guide explains how voice agents connect to the Realtime API over different transports 72:

- **Default (WebRTC)** – audio is recorded via the microphone; to use your own stream or audio element, create an `OpenAIRealtimeWebRTC` instance and pass it as the transport 73.
- **WebSocket** – for server-side use cases (e.g., phone agents) pass `transport: 'websocket'` or instantiate `OpenAIRealtimeWebSocket` 74.
- **SIP** – use `OpenAIRealtimeSIP` to bridge calls from providers like Twilio. Steps include building initial config with `buildInitialConfig()`, attaching a `RealtimeSession` with the SIP transport, and connecting with the provider's `callId` 75.
- **Cloudflare note** – Cloudflare Workers cannot open WebSocket connections directly; use `CloudflareRealtimeTransportLayer` provided by the extensions package (described later) 76.
- **Building your own transport** – implement `RealtimeTransportLayer` and emit `RealtimeTransportEventTypes` to connect with alternative speech-to-speech APIs 77.

The guide also discusses interacting directly with the Realtime API by accessing the transport layer via `session.transport` or using it standalone (thin client) 78.

---

## Extensions

### Using Any Model with the AI SDK 79

This extension integrates the **Vercel AI SDK** to use non-OpenAI models with the Agents SDK:

1. Install the extensions package: `npm install @openai/agents-extensions` 80.
2. Choose a model package from the AI SDK (e.g., `@ai-sdk/openai`) and install it 81.
3. Import the adapter and model:

```
import { openai } from '@ai-sdk/openai';
import { aisdk } from '@openai/agents-extensions/ai-sdk';
const model = aisdk(openai('gpt-5-mini'));
const agent = new Agent({ name: 'My Agent', instructions: '...', model });
run(agent, 'Question');
```

1. Provider metadata - pass provider-specific options via `providerData` (Agents SDK) which map to `providerMetadata` when using AI SDK 82.
2. UI helpers - `createAiSdkTextStreamResponse()` and `createAiSdkUiMessageStreamResponse()` wrap stream results into HTTP responses suitable for UI routes 83.

### Connect Realtime Agents to Twilio 84

- Twilio's **Media Streams API** sends raw audio from phone calls to a WebSocket server; this can drive a voice agent.
- Use the default Realtime transport in `websocket` mode, but it requires correct audio format and interruption timing.
- **TwilioRealtimeTransportLayer** (extensions package) simplifies the integration by handling connection, interruptions and audio forwarding 85.
- **Setup steps:**
  - Ensure you have a Twilio account and phone number 86.
  - Set up a WebSocket server (use ngrok or Cloudflare Tunnel to expose locally) 87.
  - Install `@openai/agents-extensions` 88.
  - Import `TwilioRealtimeTransportLayer` and create a `RealtimeAgent` & `RealtimeSession`, passing the transport 89.
  - Call `session.connect({ apiKey: ... })` to connect 90.
  - **Tips** - create the transport layer as soon as you have the WebSocket connection; call `session.connect()` immediately; listen to `transport_event` events to inspect raw Twilio messages; enable debug logs with `DEBUG=openai-agents:extensions:twilio*` 91.
  - A full example server using Fastify shows how to handle Twilio websockets and integrate tools like weather and secret functions 92.

## Realtime Agents on Cloudflare 93

- Cloudflare Workers cannot open outbound WebSockets using the global WebSocket constructor. The **CloudflareRealtimeTransportLayer** uses `fetch()` with `Upgrade: websocket` to work within workerd environments 93.
  - **Setup:**
    - Install `@openai/agents-extensions` 94.
    - Import `CloudflareRealtimeTransportLayer`, `RealtimeAgent` and `RealtimeSession` 95.
    - Create a transport specifying the realtime endpoint (e.g., `wss://api.openai.com/v1/realtime?model=gpt-realtime`) and attach it to the session 96.
    - Connect the session with an API key 97.
  - **Notes** – the Cloudflare transport uses `fetch()` with WebSocket upgrade; all `RealtimeSession` features (tools, guardrails) still work; enable debug logs for development 98.
- 

## Conclusion

The OpenAI Agents SDK for TypeScript offers a flexible framework for building both text- and voice-based agentic applications. It provides robust tooling (tools, guardrails, sessions, tracing), supports multi-agent orchestration and handoffs, and integrates with external services via MCP servers and extensions. The addition of **Voice Agents** and **Realtime transports** enables real-time speech-to-speech interactions over WebRTC, WebSocket, SIP, Twilio and Cloudflare environments, while the AI SDK adapter allows using non-OpenAI models. Guardrails, human-in-the-loop approvals and detailed tracing ensure safe, controlled operations. Overall, the SDK empowers developers to craft sophisticated agent workflows with strong observability and safety.

---

1 OpenAI Agents SDK TypeScript | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/>

2 Running agents | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/running-agents/>

3 Agents | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/agents/>

4 24 25 26 27 Guardrails | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/guardrails/>

5 19 Sessions | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/sessions/>

6 38 39 40 41 Tracing | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/tracing/>

7 48 Voice Agents | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/voice-agents/>

8 20 21 22 23 Models | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/models/>

9 15 16 Tools | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/tools/>

10 Context management | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/context/>

11 14 Results | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/results/>

12 Orchestrating multiple agents | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/multi-agent/>

13 28 29 30 Streaming | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/streaming/>

17 34 35 36 37 Model Context Protocol (MCP) | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/mcp/>

18 Handoffs | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/handoffs/>

31 32 33 Human in the loop | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/human-in-the-loop/>

42 43 44 Configuring the SDK | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/config/>

45 46 Troubleshooting | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/troubleshooting/>

47 Release process | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/release/>

49 50 51 52 53 54 Voice Agents Quickstart | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/voice-agents/quickstart/>

55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 Building Voice Agents | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/voice-agents/build/>

72 73 74 75 76 77 78 Realtime Transport Layer | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/guides/voice-agents/transport/>

79 80 81 82 83 Using any model with the Vercel's AI SDK | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/extensions/ai-sdk/>

84 85 86 87 88 89 90 91 92 Connect Realtime Agents to Twilio | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/extensions/twilio/>

93 94 95 96 97 98 Realtime Agents on Cloudflare | OpenAI Agents SDK

<https://openai.github.io/openai-agents-js/extensions/cloudflare/>